

Implementing Artificial Intelligence Features In a No-Limit Texas Hold'em Poker Bot

By Zack McGinnis

Overview

- The game of Texas Hold'em Poker
- Summary of my two added AI features:
 - SPR-Weighted UCT heuristic to the MCTS (Monte Carlo Tree Search) Bot.
 - Opponent model used to predict the probability distribution of the opponents next move.
- Results of my experiments
- Conclusions
- Demonstration?

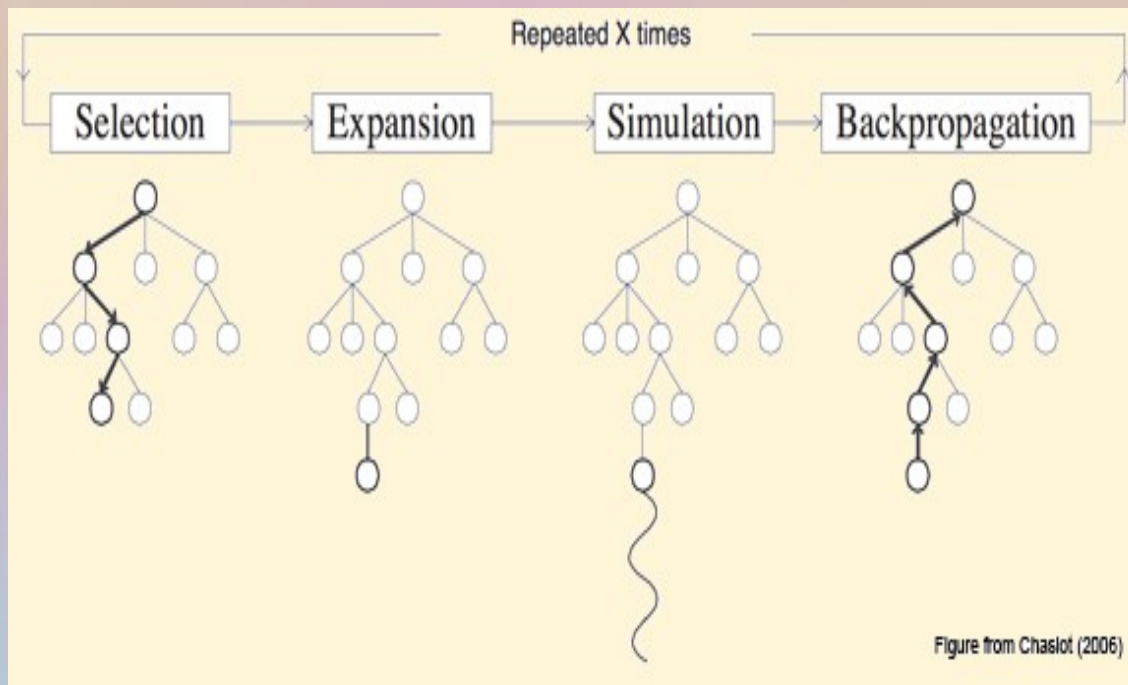
Texas Hold'Em poker: a brief overview



- Difficult to account for stochastic elements (unseen cards, bluffing, etc)
- Players have the option to: check, bet, call, raise, or fold
- My bot will be playing against one opponent
 - No-limit betting rules (not fixed limit)
 - Our opponent will be a rule-based bot, without the capability to learn from previous games.

The MCTS (Monte Carlo Tree Search) Bot

- Modified existing MCTS Bot from open-source project
 - <https://code.google.com/p/opentestbed/>
- Monte Carlo Search Tree Algorithm
 - **Selection:** Starting at root node R, recursively select optimal child nodes (explained below) until a leaf node L is reached.
 - **Expansion:** If L is not a terminal node (i.e. it does not end the game) then create one or more child nodes and select one C.
 - **Simulation:** Run a simulated playout from C until a result is achieved.
 - **Backpropagation:** Update the current move sequence with the simulation result (estimated value, and visit



Upper Confidence Bound as applied to Trees (UCT)

- Used in the first stage (selection) of MCTS
 - Acts as a heuristic to MCTS

$$\hat{V}(c_i) + C \sqrt{\frac{\ln T(P)}{T(c_i)}}$$

- Where:
 - $V(c_i)$ is the estimated value of the node,
 - (c_i) is the number of the times the node has been visited
 - P is the total number of times that its parent has been visited.
- C is a tunable bias parameter.

My weighted UCT implementation

- My UCT implementation will weight branches containing moves resulting in a higher SPR more heavily.
 - This will favor more aggressive actions throughout each iteration of the MCTS
 - e.g. More betting and raising, less checking and calling.
- SPR = Stack-to-Pot Ratio
 - ex. playerStack = \$100, currentPotSize = \$25, then SPR = 4.0
 -

$$\hat{V}(c_i) + C \sqrt{\frac{\ln T(P)}{T(c_i)}} \times (S)$$

Where

$$S = 1/\text{SPR}$$

The Opponent Model

- Utilized during the "simulation" stage of MCTS

- NaiveBayes classifier
- WEKA (Waikato Environment for Knowledge Analysis) machine learning library

The diagram shows the formula for Bayes' Theorem:
$$P(c|x) = \frac{P(x|c)P(c)}{P(x)}$$
 Arrows point from the following labels to the corresponding parts of the formula: 'Likelihood' points to $P(x|c)$; 'Class Prior Probability' points to $P(c)$; 'Posterior Probability' points to $P(c|x)$; and 'Predictor Prior Probability' points to $P(x)$.

- $P(c|x)$ is the posterior probability of class (target) given predictor (attribute)
- $P(c)$ is the prior probability of class
- $P(x|c)$ is the likelihood which is the probability of predictor given class
- $P(x)$ is the prior probability of predictor

- Trained on hand history data (10,000 hands) collected from a testbot vs. testbot game

The Opponent Model

- Models trained for many different gamestates:
 - 820 instances of post-flop check or bet actions
 - 2363 instances of post-flop fold, call, or raise actions
 - 10 instances of pre-flop check or bet actions
 - 6973 instances of pre-flop fold, call, or raise actions.
- Why NaiveBayes
 - Nominal actions of our opponent are classified based on attributes collected from previous hands
 - We can assign a probability of a class (action) occurring, given attribute values (approximately 32 per gamestate).
 - e.g. We can determine that the opponent will fold to a flop bet 80% of the time, so we assign 80% to our opponent node probability distribution for that particular gamestate

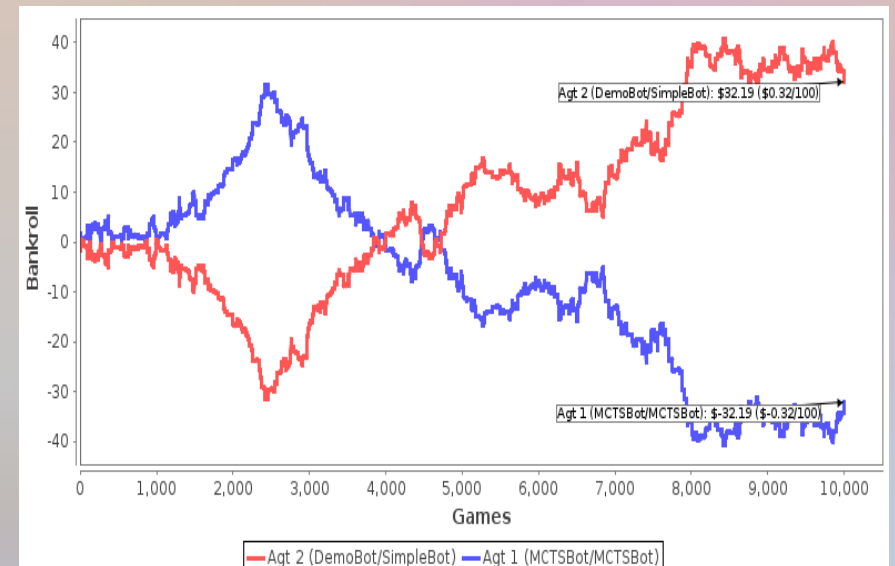

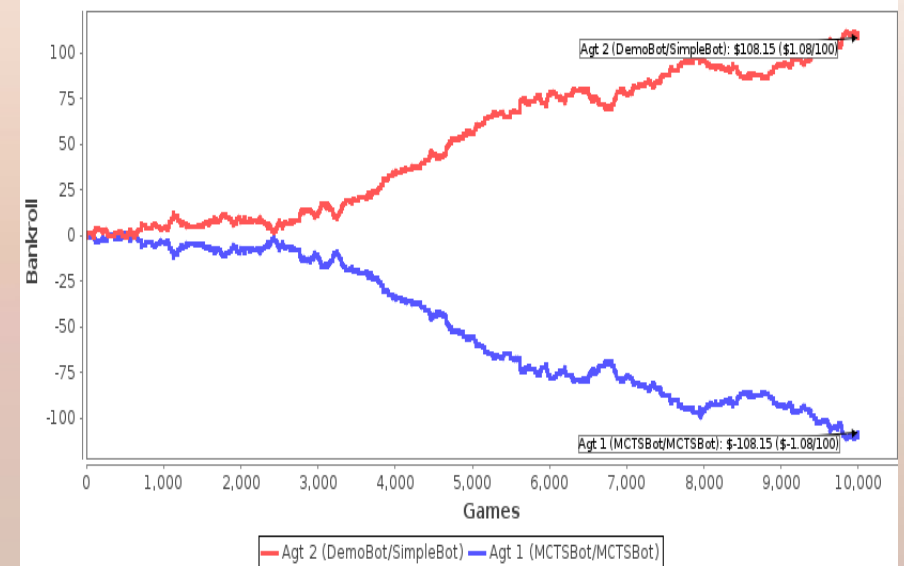

Experiment Setup

- Our opponent (testbot)
 - Rule based, (e.g. only makes decisions based on it's own two hole cards, and the exposed community cards)
 - No opponent model or ability to learn from prior games
- 10,000 games (hands) per test
- Y-axis = profit (measured in terms of bankroll (\$1.00 per bankroll))
- X-axis = # of games
- Blue line = our MCTS bot
- Red line = testbot (opponent)

Experiment Setup

- Experiment #1
 - Standard UCT with no trained opponent model
- Experiment #2
 - Standard UCT with NaiveBayes trained opponent model
- Experiment #3
 - SPR weighted UCT with no trained opponent model
- Experiment #4
 - SPR weighted UCT with NaiveBayes trained opponent model

Results



Experiment Setup

- Experiment #1

- Standard UCT with no trained opponent model
- Lost at a rate of - \$1.73/100 games

- Experiment #3

- SPR weighted UCT with no trained opponent model
- Lost at a rate of - \$0.63/100 games

- Experiment #2

- Standard UCT with NaiveBayes trained opponent model
- Lost at a rate of - \$1.08/100 games

- Experiment #4

- SPR weighted UCT with NaiveBayes trained opponent model
- Lost at a rate of - \$0.32/100 games

Conclusions

- I'm less of a loser!
- Strange activity occurring around 2500-3000 games
 - Variance?
- Would like to create and test opponent models using different classifiers (SVM, Multi-layer perceptron, etc)

Demonstration?

- MCTS bot vs testbot
 - Visualization of game-tree and MCTS algorithm iterations